

Daily Scholar Papers Report — 2026-05-11

2026-05-11

Daily Scholar Papers Report — 2026-05-11

[Download PDF](#)

Window covered: 2026-05-10 → 2026-05-11 (Google Scholar alerts + user-curated self-emails, last 24 h)

Executive Summary

An **empty alert window**. No new Scholar-alert threads landed in the last 24 hours, and no user-curated self-emails were forwarded for review either. The most recent Scholar batch resolves to 2026-05-09 16:40 UTC — ten threads already fully absorbed into the 2026-05-10 report (LiveFMBench, EvoPoC, Semantic Reification, RSafe, plus five off-topic skips). The most recent user-curated self-email is from 2026-05-09 08:08 UTC (also already absorbed, as a re-flag of arXiv 2605.06184 / SET). A broader fallback query (`after:2026/05/10`) was run on both inboxes to rule out a syntax glitch; both confirmed zero net-new threads in the 24 h window.

No Stage-1 screening or Stage-2 deep-reads were performed today. Preference posteriors were still refreshed (19 active attribute Beta-posteriors in `.private/preferences.json`) so tomorrow's run starts from the latest belief state.

Outstanding: 0 · **Keep:** 0 · **Borderline High-Priority:** 0

Highlighted Papers

No papers in window.

Cross-Paper Synthesis

Not applicable for an empty window. The standing themes from this week remain open: (i) **LLM-driving-the-search, sound non-LLM tool guarding the inner loop** — the through-line that connects Cottontail / SET (2026-05-08, 2026-05-09) with LiveFMBench's prover-checked spec-gen pipeline, EvoPoC's two-stage SMT + asset-state validator, and Semantic Reification's typed-IR-then-reify random program generator (all 2026-05-10) — versus alternative architectures that let the LLM produce the final artefact directly; and (ii) **contamination-controlled live benchmarks** — LiveFMBench's 360 post-2025 ACSL programs join REBench (2026-05-06) as templates for future evaluation design, with the open question of how durable either suite's freshness guarantee is under continued model-vendor scraping.

Writing & Rationale Insights

No new prose to mine today. A meta-observation worth recording: across the last twelve reports, the day-after-a-Saturday-/Sunday-batch window has produced zero papers four times and one paper twice. The expected daily yield of ~1.5 papers is heavily weighted toward Tue + Fri + Sun Scholar digests; an empty Sunday → Monday gap should not be treated as a pipeline bug. The infrastructure — preferences refresh, Gmail query, deploy-key push, mkdocs strict build, weasyprint render — was exercised end-to-end on this empty run and passed without incident, which is the secondary value of running the pipeline on quiet days. The cadence asymmetry also argues for keeping the report cadence at “every day” rather than “only when papers arrive”: the empty-day notes are where the cross-week themes get re-checked against memory, and skipping them silently would lose that signal.