

Daily Scholar Papers Report — 2026-05-27

2026-05-27

Daily Scholar Papers Report — 2026-05-27

[Download PDF](#)

Window covered: 2026-05-26 → 2026-05-27 (Google Scholar alerts + user-curated self-emails, last 24 h)

Executive Summary

An **empty alert window**. No new Scholar-alert threads landed in the last 24 hours, and no user-curated self-emails were forwarded for review either. The most recent Scholar batch resolves to 2026-05-25 18:29 UTC — four threads carrying seven unique papers already fully absorbed into the 2026-05-26 report (AlxCC SoK, FANDANGO, SymTEE as Outstanding; TEERepair, FuzzingBrain V2, QuartetFuzz, KnitFuzz as Keep). The most recent user-curated self-email is from 2026-05-25 10:44 UTC (re-flag of arXiv 2602.07666 / AlxCC SoK, already absorbed as the 1.1 user-pick in 2026-05-26). A broader fallback query (`after:2026/05/26`) was run on both inboxes to rule out a syntax glitch; both confirmed zero net-new threads in the 24 h window.

No Stage-1 screening or Stage-2 deep-reads were performed today. Preference posteriors were still refreshed (19 active attribute Beta-posteriors in the private store) so tomorrow's run starts from the latest belief state.

Outstanding: 0 · **Keep:** 0 · **Borderline High-Priority:** 0

Highlighted Papers

No papers in window.

Cross-Paper Synthesis

Not applicable for an empty window. The standing themes from this week remain open: (i) **LLM as a semantic proposal engine, with a deterministic backend discharging every soundness-relevant verdict** — the through-line that ties AlxCC's fuzzer ↔ LLM mutual-seeding pipelines, FANDANGO's genetic-search loop over grammar + Python semantic constraints, and SymTEE's LLM-synthesised KLEE-compatible harnesses with lightweight stubs (all 2026-05-26); (ii) **competition-driven SoKs as a methodological template** — AlxCC SoK is the immediate exemplar of "gather primary sources from every finalist immediately after the event", and the open question is whether the next major autonomous-security competition will adopt the same methodology; (iii) **TEE attack-surface tooling beyond memory-**

safety — SymTEE and its TEERepair companion extend coverage to missing-input-validation and partitioning issues, sitting alongside the older TEE-fuzzing line. None of these need fresh ink today; they re-enter scope when the next batch lands.

Writing & Rationale Insights

No new prose to mine today. A meta-observation worth recording: across the last fourteen reports, the day-after-a-Mon-evening-batch window has produced zero candidates four times now (2026-05-07, 2026-05-11, 2026-05-13, and today), so a Tue/Wed → Wed empty gap should not be treated as a pipeline bug. The infrastructure — preferences refresh, Gmail query (strict 1-day plus a broader fallback), deploy-key git pull, mkdocs strict build, and weasyprint render — was exercised end-to-end on this empty run and passed without incident, which is the secondary value of running the pipeline on quiet days. The cadence asymmetry also argues for keeping the report cadence at “every day” rather than “only when papers arrive”: empty-day notes are where the cross-week themes get re-checked against memory, and skipping them silently would erase that signal. The 2026-05-26 batch was unusually rich (three Outstanding + four Keep, all converging on the LLM-proposes / deterministic-backend-verifies architecture), so a quiet follow-up day is also a reasonable rhythm: it lets the synthesis from yesterday settle before the next wave.