

Daily Scholar Papers Report — 2026-06-28

2026-06-28

Daily Scholar Papers Report — 2026-06-28

[Download PDF](#)

Window covered: 2026-06-27 → 2026-06-28 (Google Scholar alerts + user-curated self-emails, last 24 h)

Executive Summary

A *quiet-queue* day: the 24 h window delivered two `Recommended articles` candidates from the Scholar alert feed and no user-curated self-emails. Both candidates fall inside saturated subfields — deep-learning-based source-code vulnerability detection in a mid-tier Springer journal, and LLM-driven RTSP fuzzing seed generation in MDPI *Electronics* — and neither cleared the Stage-1 bar of {top-venue, followed-researcher, or research-value × methodology-reusability ≥ High}. The user-preference posteriors (rebuilt this morning from 19 attributes) provided no $n \geq 2$ promotion or demotion, so the unbiased Stage-1 rubric governed the run end-to-end. No deep-reads today; cross-paper synthesis is held until the next non-empty day.

Outstanding: 0 · **Keep:** 0 · **Borderline High-Priority:** 0

Highlighted Papers

No papers cleared Stage-1 today.

#	Title	Authors	Venue	Link
—	(empty)	—	—	—

Triage Footnote

For traceability, the two Stage-1-skipped candidates and the rationale recorded against each:

- *Beyond the Data: Architectural Choices for Imbalanced Learning in Vulnerability Detection* — Foulefact, Djifack, Marchetto — SN Computer Science, 2026 — [Springer link](#). Skipped: saturated DL-for-vuln-detection subfield, non-top venue, no followed-researcher override.
- *LLMSGen: Generating Diverse Seeds for Multimedia Protocol Fuzzing Leveraging Large Language Models* — Wan, Zhao, Wu, Zhu — *Electronics* (MDPI) 15(12) Article 2569, 2026 — [MDPI link](#). Skipped: saturated LLM-fuzzing-seed-generation subfield (tracks ChatAFL line), narrow RTSP-only protocol scope, no formal characterisation, mid-tier non-top venue.

Cross-Paper Synthesis

Held until the next non-empty day. The pipeline only synthesises across papers that survived Stage-2 deep-read, so today's empty deep-read queue produces no synthesis content. The empty result is itself a signal: the Stage-1 rubric is doing its work — when the 24 h window contains only saturated-subfield Recommended hits at non-top venues with no followed-researcher or user-pick override, the report stays empty rather than padding with low-signal summaries.

Writing & Rationale Insights

- *Stage-1 is the rate-limiter.* On low-volume days like today the value of the digest is the *strictness* of the rubric, not the number of papers it produces. Skipping two saturated-subfield candidates is the correct output.
- *Preference signal is still thin.* Only one attribute (`venue::preprint`, $n=2$) currently has the $n \geq 2$ threshold required for promotion/demotion. As feedback accumulates, the preference layer will start overriding the unbiased rubric more often — useful to keep ingesting GitHub-issue feedback even on quiet days.
- *Window scope is intentional.* The 24 h `newer_than:1d` query is narrow enough that an empty deep-read day is expected $\approx$ once every 1–2 weeks. The next non-empty day will resume normal Highlighted Papers + Cross-Paper Synthesis structure.